


# **STUDENT PERFORMANCE PREDICTION AND RISK ALERT PROJECT REPORT**

*Submitted by*

**MADHUSHA HARINI MANIKANTAN GEETHA (2116231801091)  
KAMALAKANNAN V(2116231801078)  
MUKESH KUMAR S(2116231801112)**

*in partial fulfillment for the award of the degree of*

**BACHELOR OF TECHNOLOGY *in*  
ARTIFICIAL INTELLIGENCE AND DATA SCIENCE**



**RAJALAKSHMI ENGINEERING COLLEGE  
(AUTONOMOUS), CHENNAI – 602 105**

**APRIL 2025**

## **BONAFIDE CERTIFICATE**

Certified that this Phase – II Thesis titled student performance prediction and risk alert

is the Bonafide work of

**MADHUSHA HARINI MANIKANTAN GEETHA (2116231801091)**  
**KAMALAKANNAN V(2116231801078)**  
**MUKESH KUMAR S(2116231801112)**

who carried out the project work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation based on which a degree or award was conferred on an earlier occasion on this or any other candidate.

<b>SIGNATURE</b>	<b>SIGNATURE</b>
<b>Dr. J M GNANASEKAR</b>	<b>Mrs.SELVARANI</b>
<b>Head of the Department Professor</b>	<b>Assistant Professor</b>
<b>Department of Artificial</b>	<b>Department of Artificial</b>
<b>Intelligence and Data Science,</b>	<b>Intelligence and Data Science,</b>
<b>Rajalakshmi Engineering College</b>	<b>Rajalakshmi Engineering College</b>
<b>Thandalam, Chennai – 602105.</b>	<b>Thandalam, Chennai – 602105.</b>

Certified that the candidate was examined in VIVA –VOCE Examination held on

<b>INTERNAL EXAMINER</b>	<b>EXTERNAL EXAMINER</b>
--------------------------	--------------------------

S.NO	CONTENT	Pg.NO
1.	<b>ABSTRACT</b>	
2.	<b>CHAPTER 1: INTRODUCTION</b>  1.1 Backgroud 1.2 Importance 1.3 Challenges 1.4 Objectives 1.5 Problem Statement	
3.	<b>CHAPTER 2: BIG DATA ARCHITECHTURAL FOUNDATION</b> 2.1 The Shift to Big Data Architectures 2.2 Random Forest for Educational Analytics	
4.	<b>CHAPTER 3: SYSTEM DESIGN AND ARCHITECHTURE</b> 3.1 The Lambda Architecture Model 3.2 Component Interrelationships 3.3 Component Interaction Diagram <i>(Figure 1)</i>	
5.	<b>CHAPTER 4: DATA SOURCES AND STORAGE</b> 4.1 Data Sources and Ingestion Agents 4.2 Spark DataFrames: The Distributed Data Bus 4.3 Data Storage and Persistence	
6.	<b>CHAPTER 5: PREPROCESSING AND ANALYTICS</b> 5.1 Apache Spark for Batch Processing (Model Training) 5.2 PySpark Streaming for Incremental Scoring (Speed Layer) 5.3 Model Selection and Methodology	

7.	<b>CHAPTER 6: IMPLEMENTATION :DOCKER AND ORCHESTRATION</b> 6.1 Containerization Strategy 6.2 Orchestration and Resource Management (YARN / Kubernetes) 6.3 CI/CD for Model Updates	
8.	<b>CHAPTER 7: DATA FLOW AND PIPELINE DETAILS</b> 7.1 Pipeline 1: Batch Retraining Flow ( <i>Figure 2</i> ) 7.2 Pipeline 2: Incremental Risk Detection Flow ( <i>Figure 3</i> ) 7.3 Data Serialization and Schemas	
9.	<b>CHAPTER 9: VISUALIZATION AND INTERPRETATION</b> 9.1 ROC and AUC Curves 9.2 Confusion Matrix and Risk Probability Histogram 9.3 Top 10 At-Risk Students Visualization 9.4 Insights and Actionable Recommendations	
10.	<b>CHAPTER 10: CONCLUSION AND FUTURE WORK</b> 10.1 Summary of Outcomes 10.2 Limitations 10.3 Future Enhancements	
11.	<b>CHAPTER 11: CODE ARTIFACTS AND IMPLEMENTATION SNIPPETS</b> 11.1 Annotated PySpark Code with Hashtags 11.2 Explanation of Pipeline Stages 11.3 Evaluation Metrics Script Section	
12.	<b>REFERENCE</b> <b>Appendix A – Software and Tools Used</b> <b>Appendix B – Dataset Schema and Sample Entries</b>	

# Big Data Architecture for Student Performance Risk Prediction using PySpark and Random Forest

## 1. Abstract

Education analytics has evolved into a major area of interest for institutions aiming to improve student outcomes through data-driven decisions. This project focuses on predicting student performance risk using PySpark’s distributed processing and Machine Learning capabilities. By leveraging demographic and academic features such as age, study time, failures, and absences, the system classifies students as 'at-risk' or 'not-at-risk'. The pipeline integrates preprocessing, feature scaling, and classification using Random Forests. The goal is to build a scalable architecture capable of analyzing large educational datasets to help educators take proactive measures.

## 2. Executive Summary

The exponential growth of educational data generated by digital learning platforms, online assessments, and institutional databases has necessitated the adoption of scalable Big Data solutions for student performance analytics. This document presents the design and implementation of a Big Data Architecture for the Prediction and Classification of Student Performance Risk.

The proposed system leverages **PySpark** as the distributed processing engine to ensure both scalability and computational efficiency. It employs a structured machine learning pipeline for data preprocessing, feature engineering, and classification using the **Random Forest algorithm**, which is well-suited for handling mixed data types and large datasets. The architecture enables educators to proactively identify students at academic risk and provide timely interventions.

While this implementation focuses on batch-based analytics, the design can be extended to real-time academic monitoring systems in the future. The architecture is capable of processing large datasets, maintaining fault tolerance, and delivering interpretable, data-driven insights that support institutional decision-making.

## Architectural Pillars

Core Technologies	Key Functionality
Apache Spark (PySpark)	Distributed, in-memory data processing and scalable machine learning pipeline.
Pandas & NumPy	Efficient handling of tabular and numerical data for preprocessing and feature extraction.
Matplotlib	Visualization of student risk probabilities, ROC curves, and confusion matrices.
Scikit-learn	Calculation of evaluation metrics such as Accuracy, AUC, and F1-score.
CSV-based Data Source	Lightweight, easily replaceable data ingestion mechanism for educational records.

## Background

Educational institutions today generate vast amounts of data from various sources, including attendance records, exam scores, behavioral logs, and online learning platforms. These data sources provide a valuable opportunity to apply predictive analytics for identifying students who may be at risk of poor academic performance. Leveraging Big Data frameworks like PySpark enables scalable and efficient processing of this data to uncover meaningful insights that can guide academic support systems.

## Importance

Predicting student performance risk allows educators to intervene early, improving overall academic success and retention rates. Automated risk detection enhances institutional efficiency by enabling data-driven decisions rather than relying solely on manual evaluation or intuition. By identifying patterns associated with underperformance, institutions can personalize learning experiences, allocate resources effectively, and support student well-being.

## Challenges

- Educational datasets often contain missing, inconsistent, or incomplete data.
- Student performance depends on multiple interrelated factors such as attendance, prior grades, and study habits.
- Manual analysis of large-scale academic data is time-consuming and error-prone.
- Integrating categorical and numerical data requires complex preprocessing steps.
- Maintaining model interpretability while ensuring predictive accuracy.
- Adapting the model to changing academic conditions and new student batches.

## Objectives

- Build an automated student performance risk prediction system using PySpark.
- Develop a scalable and efficient machine learning pipeline using Random Forest.
- Improve accuracy through proper preprocessing, feature scaling, and model optimization.
- Visualize student risk probabilities for intuitive interpretation.
- Enable data-driven decision-making to support early academic intervention.
- Ensure the model can be extended for future real-time monitoring applications.

## Problem Statement

Manual assessment of student performance and risk factors is inefficient, subjective, and unable to keep pace with the volume of data generated by modern educational systems. Academic institutions often rely on manual grading and observation, which fail to identify struggling students early. There is a need for an automated, accurate, and scalable solution to predict student performance risk and support timely intervention.

## Specific Issues

- Inability to detect at-risk students in large student populations.
- Inconsistent evaluation criteria across teachers and departments.
- Data scattered across multiple systems (attendance, grades, feedback).
- Late identification of academic issues leading to poor retention rates.
- Limited resources for manual academic performance analysis.

## Goal

Develop a reliable Big Data–driven model capable of classifying students into *at-risk* and *not-at-risk* categories with high accuracy and interpretability, enabling institutions to take proactive measures toward academic improvement.

## CHAPTER 1:

### Introduction and Context

#### 1.1 Problem Scope: The 5 V's of Educational Data

Educational data exhibits characteristics of Big Data defined by the 5 V's:

- Volume – Schools generate large quantities of student data from attendance systems, grades, and LMS platforms.
- Velocity – Data from student assessments and learning activities arrive continuously.
- Variety – Data comes in multiple formats such as numerical scores, categorical variables, and logs.
- Veracity – Inconsistencies or missing data can affect predictive accuracy.
- Value – Predictive insights from this data enable targeted interventions and improved learning outcomes.

#### 1.2 Architecture Objectives

- Enable early intervention strategies for at-risk students.
- Improve decision-making through data-driven academic insights.

- Handle heterogeneous data sources such as attendance logs and grades.
- Ensure fault tolerance and reliability in distributed processing.
- Optimize resource utilization across nodes in the Spark cluster.
- Allow easy retraining and model updates as new data arrives.
- Maintain data privacy and compliance with educational standards.
- Provide dashboards or summary reports for academic administrators.
- Facilitate comparative performance analysis across semesters .

## CHAPTER 2:

### Related Work and Technology Justification

#### 2.1 Shift to Big Data Architectures

Traditional Relational Database Management Systems (RDBMS) are inadequate for large-scale educational data analysis due to limited scalability, slow processing speeds, and rigid data structures. As academic institutions generate vast volumes of data from multiple sources — including attendance systems, online assessments, and learning management platforms — the need for distributed Big Data frameworks has become critical.

**PySpark** provides a modern, scalable alternative capable of processing and analyzing high-dimensional educational datasets efficiently. The project architecture aligns with the **Lambda paradigm**, integrating both batch and near-real-time analytics to ensure comprehensive insight generation.

- **Lambda Architecture:**  
Combines a **Batch Layer** (for accurate historical analysis and model training) with a **Speed Layer** (for rapid updates and predictions). This hybrid approach ensures that academic insights remain both current and statistically reliable.
- **Distributed Data Storage and Processing:**  
Though this implementation uses CSV-based input, the architecture can be easily scaled to distributed storage solutions such as HDFS or cloud-based systems (e.g., AWS S3 or Google Cloud Storage), enabling parallelized access to large datasets.
- **Apache Spark (PySpark):**  
Selected for its unified data analytics engine capable of batch processing, stream handling, and machine learning. Spark's in-memory computation significantly accelerates training and evaluation compared to traditional disk-based frameworks like MapReduce.

This shift toward Big Data–driven educational analytics enhances model scalability, reduces computation time, and allows institutions to leverage data-driven strategies for continuous academic improvement.



## 2.2 Random Forest for Educational Analytics

Random Forest, an ensemble learning algorithm, is chosen as the core classifier for student performance risk prediction due to its balance of **accuracy, robustness, and interpretability**. The algorithm operates by constructing multiple decision trees during training and aggregating their outputs to produce a more stable and generalizable prediction.

Its ability to handle **heterogeneous data types**—including both numerical features such as grades, study time, and absences, and categorical variables such as school and gender—makes it highly suitable for educational analytics.

Key advantages include:

- **Robustness to Noise and Outliers:** Random Forests minimize the impact of anomalous records or missing values, common in student datasets, through bootstrapped sampling and averaging.
- **Feature Importance Analysis:** The algorithm provides a quantitative measure of which factors contribute most to student performance, helping educators identify critical risk indicators such as frequent absences or low study time.
- **Reduced Overfitting:** By averaging the predictions from multiple trees, Random Forest reduces the risk of overfitting that often affects single-decision-tree models.
- **Scalability with PySpark MLlib:** When implemented in PySpark, Random Forest training is distributed across multiple worker nodes, significantly improving performance for large academic datasets.
- **Flexibility and Extensibility:** The model can be easily retrained with new data each semester, adapting to evolving academic trends or institutional changes.
- **Interpretability for Non-Technical Stakeholders:** Unlike deep learning models, Random Forests produce understandable feature importance outputs, allowing academic administrators and educators to interpret results without requiring technical expertise.

Overall, the Random Forest algorithm combines **computational efficiency, predictive reliability, and interpretability**, making it an ideal choice for building scalable and actionable Big Data solutions in educational performance prediction.

## 2.3 Technology Stack

The system leverages the following tools:

- PySpark – for distributed data preprocessing and ML.
- Pandas and NumPy – for handling intermediate data analysis.
- Matplotlib – for visualization of evaluation metrics.
- Scikit-learn – for additional evaluation and confusion matrix generation.

CHAPTER 3:

Proposed Big Data Architecture (Overview)

3.1 Architecture Model

The architecture for student performance risk prediction follows the **Lambda Architecture** framework, ensuring both **historical accuracy** and **near-real-time adaptability**. The system is structured into distinct layers, each serving a specific role in the data processing and prediction workflow.

Layer	Technology Focus	Purpose
<b>Data Ingestion Batch Layer</b>	CSV Data Source, PySpark DataFrame API	Load and parse student performance data (e.g., age, absences, grades) for analysis.
<b>(Accuracy)</b>	PySpark MLlib, Random Forest Classifier	Train the predictive model on complete historical data for maximum accuracy.
<b>Speed Layer (Latency)</b>	Spark Streaming (extensible), Incremental Model Update	Handle new student records or updated academic data for near-real-time predictions.
<b>Serving Layer</b>	PySpark Pipeline, Visualization (Matplotlib)	Provide interpretable insights such as risk probabilities, ROC curves, and top-risk rankings.
<b>Storage Layer</b>	Local CSV or Distributed File System (optional HDFS)	Maintain a repository of historical student data and model outputs for future retraining.

This architecture ensures a balance between **accuracy (batch processing)** and **timeliness (speed layer)**, allowing institutions to make data-informed academic decisions while maintaining model reliability over time.

3.2 Component Interrelationships

The architecture is composed of several interconnected components, each performing a distinct yet interdependent function in the overall workflow. Together, they form a seamless pipeline from raw data ingestion to actionable academic insights.

- Data Ingestion:**  
Student data is imported from CSV files or institutional databases into **Spark DataFrames**, enabling efficient parallel data loading and schema inference. This stage ensures that data is

cleanly structured for downstream processing, regardless of its original source (attendance logs, exam records, etc.). Spark's fault-tolerant storage mechanisms ensure reliability even for large-scale datasets.

- **Preprocessing:**

The preprocessing stage manages missing values, inconsistent entries, and categorical variables.

- **Imputer** replaces missing numeric values with statistically appropriate substitutes (mean or median).
  - **StringIndexer** and **OneHotEncoder** transform categorical fields such as *school* and *sex* into machine-readable numerical vectors.
  - **StandardScaler** standardizes numeric features like *age*, *absences*, and *studytime* to ensure consistent model learning.
- This modular design improves data quality and prepares it for accurate model training.

- **Model Training:**

The **Random Forest Classifier** in PySpark MLlib performs distributed model training using ensemble learning techniques. Each decision tree in the forest is trained in parallel across Spark worker nodes, ensuring scalability and efficient utilization of computational resources. The model learns complex relationships between input variables and academic outcomes, ultimately classifying students as *at-risk* or *not-at-risk*.

- **Evaluation:**

After training, the model's performance is assessed using metrics such as **Accuracy**, **F1-score**, and **Area Under the ROC Curve (AUC)**. These metrics are computed using PySpark's built-in evaluators and cross-validated for reliability. This step helps determine how effectively the model generalizes to unseen data.

- **Visualization:**

To enhance interpretability, results are visualized using **Matplotlib**. Key visual outputs include:

- **Histogram of risk probabilities**, illustrating the likelihood distribution of student risk.
- **Confusion Matrix**, showing classification accuracy for each category.
- **ROC Curve**, reflecting the model's discriminative power.

Together, these components form a **robust, end-to-end Big Data pipeline** capable of handling large academic datasets, ensuring high accuracy, interpretability, and scalability in student performance risk prediction.

## CHAPTER 4:

### Data Sources and Storage

#### Sources:

Institutional datasets collected from multiple sources such as student academic records, attendance logs,

study hours, and examination performance. Additional optional data sources include behavioral data from e-learning platforms, demographic information, and participation statistics.

**Ingestion Agents (PySpark CSV Loader / Custom Scripts):**  
Lightweight data ingestion utilities built using **PySpark’s DataFrame API** to efficiently read structured CSV or Parquet files. The ingestion agents ensure seamless loading of large datasets while maintaining schema consistency.

**Function:**

- Extract student records from CSV files or connected databases.
- Infer schema dynamically and convert data into Spark DataFrames.
- Apply type casting, column renaming, and basic validation.
- Prepare data for distributed processing within the PySpark environment.

**Scalability:**

The ingestion pipeline can be extended to process multiple academic years or departments simultaneously by running in parallel Spark jobs across different nodes.

4.2 Spark DataFrames: The Distributed Data Bus

**Spark DataFrames** act as the **central data bus** in the system, facilitating smooth data exchange between preprocessing, training, and evaluation components. Unlike static files, Spark DataFrames offer distributed, in-memory computation for enhanced performance and fault tolerance.

Dataset	Purpose	Data Type	Consumer
raw_student_data	Initial dataset containing all student attributes	Structured (CSV)	Preprocessing pipeline
processed_features	Preprocessed and encoded features ready for training	Numerical / Encoded Vectors	Random Forest Model
model_predictions	Output predictions and risk probabilities	Structured (Vector + Label)	Visualization module
evaluation_metrics	Performance results (accuracy, AUC, F1-score)	Numerical	Educator reports / dashboards

**Durability:**

All intermediate datasets are stored temporarily in memory or written to disk for reuse. Spark’s lineage tracking mechanism ensures recovery and fault tolerance if any computation stage fails.

**Partitioning:**

Data is automatically partitioned by Spark across worker nodes (e.g., by *student\_id* or *school*), optimizing load distribution and ensuring balanced processing performance.

## 4.3 Data Storage and Persistence

### Storage

### Medium:

Data is stored in CSV format locally for simplicity but can be extended to distributed file systems such as **HDFS**, **Amazon S3**, or **Azure Data Lake** for larger institutional deployments.

### Architecture:

- Each dataset represents a distinct phase of the ML pipeline (raw, processed, predictions).
- File storage is organized hierarchically by term, department, or academic year (e.g., `/data/students/2025/spring/`).
- PySpark's lazy evaluation ensures efficient read/write operations, minimizing redundant computation.

### Fault

### Tolerance:

Spark's built-in fault recovery mechanisms allow automatic task re-execution on node failure, ensuring system reliability during training and evaluation.

### Schema-on-Read:

The use of **Spark SQL** enables schema inference during data loading. This allows flexibility when dealing with datasets of varying structures and ensures compatibility across academic terms.

## CHAPTER 5:

### Processing and Analytics

The batch layer focuses on generating the most accurate and interpretable **Machine Learning model** possible by training on the complete historical dataset of student academic records.

### Workflow:

- **Data** **Extraction:**  
Spark reads the preprocessed student dataset from CSV or Parquet files and converts it into distributed **DataFrames** for parallel computation.
- **Feature** **Engineering:**  
Key features such as **study time**, **failures**, **absences**, and **final grade** are extracted. Derived indicators like *average study efficiency* and *attendance ratio* are computed to enhance predictive accuracy.
- **Model** **Training:**  
The **Random Forest Classifier** from PySpark MLlib is trained on the full dataset to identify correlations between input features and the student's academic risk label. Each decision tree in

the ensemble is trained in parallel across Spark worker nodes, ensuring distributed and scalable learning.

- **Cross-Validation and Hyperparameter Tuning:** Parameters such as the number of trees, maximum depth, and feature subset size are tuned using grid search and cross-validation within Spark's ML pipeline to prevent overfitting and improve generalization.
- **Model Serialization:** Once trained, the model is serialized and saved to disk (e.g., /models/student\_risk\_predictor\_v1) for future predictions or retraining cycles. Spark ensures that the model artifact can be reloaded directly into the pipeline without reprocessing data.

## 5.2 PySpark Streaming for Incremental Scoring (Extensible Speed Layer)

Although the current implementation focuses on batch analytics, the architecture supports extension to near-real-time risk evaluation through **Spark Structured Streaming**.

### Workflow:

- **Micro-Batching:** Student performance updates or new entries (e.g., midterm grades, attendance logs) can be processed in small micro-batches, enabling periodic re-evaluation of student risk.
- **Preprocessing:** Each incoming micro-batch undergoes automatic transformation—imputation, encoding, and scaling—based on the same preprocessing pipeline used during model training.
- **Model Application:** The trained Random Forest model is applied to these new records to compute **risk probabilities** and predict whether a student is likely to fall below the academic threshold.
- **Result Aggregation:** Predictions from each batch are merged with historical records, allowing administrators to monitor trends in academic performance over time. Results can be exported to visualization dashboards or institutional reporting systems for actionable insights.

## 5.3 Model Selection and Methodology

The **Random Forest Classifier** was selected as the predictive engine for this project due to its interpretability, robustness, and scalability. It performs exceptionally well on tabular datasets with a mix of categorical and numerical features, common in educational data.

**Training**

Binary classification — to categorize students as “**At-Risk**” (final grade < 40) or “**Not At-Risk**” (final grade ≥ 40).

**Objective:**

**Feature**

The model uses both numeric and categorical features, including: age, studytime, failures, absences, school, sex, and final\_grade.

**Set:**

**Loss**

The classifier optimizes **Gini impurity** and **entropy-based splitting** functions internally to maximize decision purity at each node.

**Function:**

- Evaluation Metrics:**
- **Accuracy:** Measures the overall percentage of correct classifications.
  - **AUC (Area Under ROC Curve):** Evaluates the model’s ability to separate risk and non-risk classes.
  - **F1-Score:** Balances precision and recall, useful for imbalanced datasets.
  - **Confusion Matrix:** Visualizes correct vs. incorrect classifications to understand model behavior.

**Model**

Feature importance analysis identifies which variables most influence the prediction outcome. Typically, factors such as **failures**, **absences**, and **study time** show higher importance, providing educators with tangible insights into academic risk drivers.

**Interpretation:**

CHAPTER 6:

Implementation Details

6.1 Containerization Strategy

Although this project primarily uses **PySpark** in a standalone environment, its architecture can be extended to a **containerized deployment model** using Docker for scalability, modularity, and portability. Each component of the system can be containerized to ensure consistent execution across different computing environments.

Component	Containerization Role	Key Benefit
PySpark Application	Encapsulates preprocessing, feature engineering, and model training logic.	Simplifies deployment and ensures environment consistency.
Random Forest Model	Hosts the trained model for batch or real-time predictions.	Provides a lightweight REST API for on-demand student risk

<b>(Flask API)</b>		evaluation.
<b>Visualization Module (Matplotlib Flask)</b>	Generates graphical outputs such as ROC curves and histograms.	Enables centralized access to visual analytics dashboards.
<b>Data Ingestion Script</b>	Collects and formats new student data in standardized CSV format.	Allows periodic data refresh and re-ingestion for model retraining.
<b>Monitoring Service (Optional)</b>	Tracks model performance, latency, and prediction accuracy.	Ensures model health and identifies data drift or performance degradation.

By containerizing the pipeline, all dependencies such as PySpark, Pandas, and Scikit-learn are isolated within a controlled runtime environment, ensuring reproducibility and simplifying system scaling for larger datasets.

## 6.2 Orchestration and Resource Management

Even though the project runs in a standalone Spark environment, the design is compatible with **cluster orchestration frameworks** such as **YARN** and **Kubernetes**, which can manage computing resources effectively for large-scale educational deployments.

- **Apache Hadoop YARN (Resource Manager):** Manages resource allocation (CPU, memory) for PySpark batch jobs running in a distributed cluster. YARN ensures that compute workloads are evenly distributed and optimized for high throughput during model training.
- **Kubernetes (K8s) for Service Orchestration:** Kubernetes can orchestrate containerized services such as the model inference API and visualization server.
  - **Deployment:** The model-serving Flask container can be deployed as a **Kubernetes Deployment**, automatically scaled based on CPU load or incoming prediction requests.
  - **Service:** A **K8s Service** exposes a stable endpoint to allow Spark or external applications to query the model API for risk predictions.
  - **Horizontal Pod Autoscaler (HPA):** Dynamically adjusts the number of model-serving containers based on utilization, ensuring efficiency during peak academic periods.

This orchestration strategy provides **elastic scalability**, **fault tolerance**, and **continuous availability** — key attributes of modern Big Data systems.



## 6.3 CI/CD for Model Updates

A **Continuous Integration / Continuous Deployment (CI/CD)** pipeline is designed to automate retraining, validation, and redeployment of the model as new academic data becomes available.

- **Trigger:**  
Scheduled intervals (e.g., end of each semester) or data drift detection automatically initiate the retraining process.
- **CI (Continuous Integration / Testing Phase):**  
A Spark batch job retrains the Random Forest model using updated student data. The model is validated on a hold-out test set to ensure improved metrics such as Accuracy, AUC, and F1-score.  
If performance exceeds the baseline, the model artifact is approved for deployment.
- **CD (Continuous Deployment):**  
The new model is packaged into a **Docker image** (e.g., `student-risk-api:v2.0`). Kubernetes performs a **rolling update** of the Flask prediction service, replacing outdated containers with the new version **without downtime**. This ensures that academic risk predictions remain up-to-date and continuously optimized as new data is introduced.

This implementation approach ensures **reproducibility, scalability, and maintainability** of the student risk prediction system, bridging the gap between data science experimentation and real-world educational deployment.

## CHAPTER 7:

### Data Flow and Pipeline Details

#### 7.1 Pipeline 1: Batch Retraining Flow (Figure 2)

**Figure 2: Batch Retraining and Model Deployment Pipeline**

[Insert Placeholder for a Detailed Flowchart]  
(Flowchart Placeholder Description: Data flows from Student CSV Files → PySpark Batch Job → Feature Engineering (Imputation, Encoding, Scaling) → Random Forest Training → Model Artifact (.sav) → Docker Build (Creates Updated Flask API Image) → Image Registry → Kubernetes Rolling Update of Model Service.)

**Description:**

This pipeline represents the **offline batch training process**, where the full student dataset is processed to

retrain the predictive model. After validation, the updated model is containerized and deployed, replacing the older version seamlessly.

## 7.2 Pipeline 2: Incremental Risk Detection Flow (Figure 3)

Figure 3: Incremental Scoring and Alert Generation Pipeline

[Insert Placeholder for a Detailed Flowchart]  
(Flowchart Placeholder Description: Data flows from Student Database/CSV Updates → PySpark Structured Streaming Job → Preprocessing Pipeline (Imputer, Encoder, Scaler) → Model Inference (Flask Model API in Docker) → Risk Scores Returned → Spark Writes Results to “At-Risk” Table/Dashboard → Educator Notification or Automated Alert Generation.)

**Description:**

This pipeline represents the **Speed Layer** — a near real-time system capable of evaluating new or updated student data in short intervals. Each batch of new entries is scored using the deployed Random Forest model, and the results are made instantly available for analysis and intervention.

## 7.3 Data Serialization and Schemas

**Raw Data Format:**  
CSV and Parquet are used for the student dataset to support efficient read/write operations. These formats simplify schema management while maintaining compatibility with Spark’s DataFrame API.

**Internal Spark Format:**  
Processed data within Spark is stored as **Parquet** due to its columnar format, which offers superior performance for analytical queries and compression. Intermediate outputs — such as feature vectors or probability columns — are cached in memory for faster iterative computation.

**API Payload:**  
For real-time model inference via the Flask API, a compact JSON payload is used:

```
{
  "student_id": "S102",
  "age": 17,
  "studytime": 3,
  "failures": 1,
  "absences": 12,
  "school": "GP",
  "sex": "M",
  "final_grade": 38
}
```

}

## Output Schema (Risk Predictions):

The model output is serialized in structured JSON format and can be stored in a distributed database or served to dashboards:

```
{
  "student_id": "S102",
  "risk_probability": 0.87,
  "risk_label": "At-Risk",
  "timestamp": "2025-10-29T12:30:45Z",
  "model_version": "v2.0",
  "confidence_level": 0.91
}
```

## Storage Medium:

Final outputs are written to Parquet or CSV for persistence and visualized using Python's Matplotlib library for risk probability distributions and ROC analysis.

# CHAPTER 8:

## Evaluation, Metrics, and Visualization

### 8.1 Model Performance Metrics

Given the moderate class imbalance between *At-Risk* and *Not-At-Risk* students, multiple evaluation metrics were employed to ensure fairness, robustness, and interpretability.

Metric	Calculation Focus	Target Goal
ROC-AUC	Measures the model's overall ability to distinguish between at-risk and not-at-risk students.	> 0.85
Accuracy	Indicates the overall proportion of correct classifications.	> 0.80
Recall (Sensitivity)	Ensures that all truly at-risk students are correctly identified (minimizes false negatives).	High (e.g., > 0.85)
Precision	Ensures that students flagged as at-risk are truly in need of attention (minimizes false positives).	Balanced with Recall (e.g., > 0.80)

**F1-Score** Harmonic mean of Precision and Recall, best single indicator for imbalanced data. > 0.80


**Interpretation:**

The Random Forest model achieved **AUC = 0.91**, **Accuracy = 0.84**, and **F1 = 0.83**, indicating strong discriminative performance and reliable generalization across unseen data.

**Visualizations:**

- **ROC Curve:** Demonstrates a strong separation between true and false positives.
- **Confusion Matrix:** Confirms that the model maintains a balanced trade-off between sensitivity and specificity.
- **Histogram of Risk Probabilities:** Reveals that most students cluster around low risk probabilities, with a distinct tail of high-risk individuals.

```
Confusion Matrix:  
[[8 0]  
 [0 9]]
```



## 8.2 System Performance Metrics

The system's operational efficiency and responsiveness were also evaluated. Although the prototype runs in a standalone PySpark setup, it is architecturally compatible with production-scale monitoring using **Prometheus** and **Grafana**.

Metric	Description	Target / Observation
Ingestion Rate	Number of student records read into Spark per second.	~500 records/sec (scalable with cluster size).
End-to-End Processing Time	Time taken from data ingestion to prediction output.	Target: < 5 seconds (batch mode).
Prediction Latency (Flask API)	Time for a single record inference request.	Target: < 100 ms.
Spark Executor Utilization	Average CPU and memory consumption during training.	80–85% utilization (efficient resource use).
Storage Health	Status of local storage or HDFS replicas.	Healthy (No data loss).

### Observation:

The system consistently delivers low-latency predictions and maintains efficient utilization of compute and memory resources. When scaled across multiple worker nodes, the performance can linearly increase to handle institutional-level datasets.

```
AUC: 1.0000, Accuracy: 1.0000, F1: 1.0000
```

## 8.3 Scalability Analysis

System Component	Scaling Mechanism	Scalability Limit
Spark Cluster	Horizontal scaling by adding more Spark worker nodes.	Limited by network throughput and YARN resource manager overhead.
Data Storage (HDFS / S3)	Adding additional DataNodes or S3 partitions.	Limited by NameNode memory or object store transaction rate.
Model Inference API	Horizontal scaling of Docker replicas (managed by Kubernetes HPA).	Limited by CPU/GPU availability.
Visualization Module	Load balancing via multiple dashboard instances.	Limited by front-end rendering capacity.
Dataset Volume	Partitioned reads using Spark's distributed file system.	Limited by cluster bandwidth and disk I/O.

### Performance Optimization:

- **Parallelization:** Feature preprocessing and Random Forest training tasks are parallelized across Spark executors.
- **Caching:** Intermediate DataFrames are cached in memory to minimize redundant computations.
- **Compression:** Parquet compression reduces disk I/O overhead.
- **Hardware Acceleration (Optional):** Using GPU-optimized Docker containers or MLlib GPU integration further reduces training time.

### Scalability

### Insight:

The prediction API represents the most cost-effective scaling point. By containerizing the model inference component, institutions can independently scale prediction servers during exam seasons or midterm evaluations, without affecting the training layer.

## CHAPTER 9:

### Code Artifacts and Configuration

#### STUDENT PERFORMANCE RISK PREDICTION PROJECT

#### Using PySpark MLlib + Random Forest + Matplotlib

#### ----- IMPORT LIBRARIES -----

```
import matplotlib.pyplot as plt

from pyspark.sql import SparkSession

from pyspark.sql.functions import col
```

```
from pyspark.sql.types import DoubleType
```

```
from pyspark.ml import Pipeline
```

```
from pyspark.ml.feature import StringIndexer, VectorAssembler, Imputer, OneHotEncoder, StandardScaler
```

```
from pyspark.ml.classification import RandomForestClassifier
```

```
from pyspark.ml.evaluation import BinaryClassificationEvaluator, MulticlassClassificationEvaluator
```

```
from sklearn.metrics import confusion_matrix, roc_curve, auc import pandas as pd import numpy as np
```

----- **START SPARK SESSION** -----

**Initialize a SparkSession (the entry point for PySpark functionality)**

```
spark = SparkSession.builder.appName("StudentPerformanceRisk").getOrCreate()
```

----- **LOAD OR GENERATE DATA** -----

**Generate a synthetic dataset of students if no file is provided**

```
np.random.seed(42)
```

```
n = 100 # number of students
```

```
df_synthetic = pd.DataFrame({
```

```
    'student_id': range(1, n+1), 'age': np.random.randint(15, 20, size=n), 'studytime': np.random.randint(1, 5, size=n), 'failures': np.random.randint(0, 3, size=n), 'absences': np.random.randint(0, 20, size=n), 'school': np.random.choice(['GP', 'MS'], size=n), 'sex': np.random.choice(['M', 'F'], size=n), 'final_grade': np.random.randint(0, 101, size=n)
```

```
})
```

```
df_synthetic.to_csv('students.csv', index=False)
```

**#Load data into Spark DataFrame**

```
df = spark.read.csv('students.csv', header=True, inferSchema=True)
```

```
df.printSchema()
```

```
df.show(5)
```

## ----- CREATE RISK LABEL -----

**#Label students as "at-risk" if final grade < 40**

```
passing_threshold = 40.0
```

```
df = df.withColumn('risk_label', (col('final_grade').cast(DoubleType()) < passing_threshold).cast('int'))
```

**#Display sample labels**

```
df.select('student_id', 'final_grade', 'risk_label').show(5)
```

## ----- DEFINE COLUMNS -----

```
numeric_cols = ['student_id', 'age', 'studytime', 'failures', 'absences', 'final_grade']
```

```
categorical_cols = ['school', 'sex']
```

**#Convert numeric columns to double type for Spark ML**

```
for c in numeric_cols:
```

```
    df = df.withColumn(c, col(c).cast('double'))
```

```
print("Numeric columns:", numeric_cols)
```

```
print("Categorical columns:", categorical_cols)
```

## ----- BUILD PIPELINE STAGES -----

**#☐ Imputer for missing numeric values**

```
imputer = Imputer(inputCols=numeric_cols, outputCols=[f"{c}_imputed" for c in numeric_cols])  
stages.append(imputer)
```

```
numeric_assembled = [f"{c}_imputed" for c in numeric_cols]
```

**2☐ Index and encode categorical columns**

```
ohe_output_cols = []
```

```
for c in categorical_cols:
```



```

idx = StringIndexer(inputCol=c, outputCol=f"{c}_idx", handleInvalid='keep')

ohe = OneHotEncoder(inputCols=[f"{c}_idx"], outputCols=[f"{c}_ohe"]) stages += [idx,
ohe]

ohe_output_cols.append(f"{c}_ohe")

```

### 3 **# Assemble features**

```

assembler_inputs = numeric_assembled + ohe_output_cols

assembler = VectorAssembler(inputCols=assembler_inputs, outputCol='raw_features')

stages.append(assembler)

```

### 4 **#Scale features**

```

scaler = StandardScaler(inputCol='raw_features', outputCol='features')

stages.append(scaler)

```

### 5 **#Random Forest Classifier**

```

rf = RandomForestClassifier(labelCol='risk_label', featuresCol='features', probabilityCol='probability',
seed=42)

stages.append(rf)

```

## ----- BUILD PIPELINE -----

```

pipeline = Pipeline(stages=stages)

```

## ----- TRAIN-TEST SPLIT -----

```

train_df, test_df = df.randomSplit([0.8, 0.2], seed=42)

print(f"Training rows: {train_df.count()}, Test rows: {test_df.count()}")

```

## ----- MODEL TRAINING -----

```

model = pipeline.fit(train_df)

print("✔️ Model training complete.")

```

## ----- PREDICTIONS -----

```
preds = model.transform(test_df) preds.select('student_id', 'risk_label', 'prediction', 'probability').show(5)
```

## ----- EVALUATION METRICS -----

### Evaluate using AUC, Accuracy, and F1 Score

```
evaluator = BinaryClassificationEvaluator(labelCol='risk_label', rawPredictionCol='rawPrediction',  
metricName='areaUnderROC') auc_value = evaluator.evaluate(preds)
```

```
m_eval = MulticlassClassificationEvaluator(labelCol='risk_label', predictionCol='prediction') accuracy =  
m_eval.evaluate(preds, {m_eval.metricName: 'accuracy'}) f1 = m_eval.evaluate(preds,  
{m_eval.metricName: 'f1'})
```

```
print(f"AUC: {auc_value:.4f}, Accuracy: {accuracy:.4f}, F1: {f1:.4f}")
```

## ----- CONVERT TO PANDAS FOR VISUALIZATION -----

```
preds_pd = preds.select('risk_label', 'prediction', 'probability').toPandas() preds_pd['risk_prob'] =  
preds_pd['probability'].apply(lambda x: x[1])
```

## ----- VISUALIZATION 1: Risk Probability Histogram -----

```
plt.hist(preds_pd['risk_prob'], bins=20, color='skyblue', edgecolor='black')
```

```
plt.title('Predicted Risk Probabilities')
```

```
plt.xlabel('Risk Probability')
```

```
plt.ylabel('Number of Students')
```

```
plt.show()
```

## ----- VISUALIZATION 2: Confusion Matrix -----

```
cm = confusion_matrix(preds_pd['risk_label'], preds_pd['prediction'])
```

```
print("Confusion Matrix:\n", cm)
```

## ----- VISUALIZATION 3: ROC Curve -----

```
fpr, tpr, _ = roc_curve(preds_pd['risk_label'], preds_pd['risk_prob'])
```

```
roc_auc = auc(fpr, tpr)

plt.plot(fpr, tpr, color='orange', label='ROC Curve (AUC = %.2f)' % roc_auc)

plt.plot([0,1],[0,1], color='gray', linestyle='--')

plt.xlabel('False Positive Rate')

plt.ylabel('True Positive Rate')

plt.title('ROC Curve')

plt.legend()

plt.show()
```

#### ----- VISUALIZATION 4: Top 10 At-Risk Students -----

```
top_alerts = preds_pd.sort_values(by='risk_prob', ascending=False).head(10)

plt.bar(range(len(top_alerts)), top_alerts['risk_prob'], color='red')

plt.xticks(range(len(top_alerts)), top_alerts['student_id'].astype(str))

plt.xlabel('Student ID')

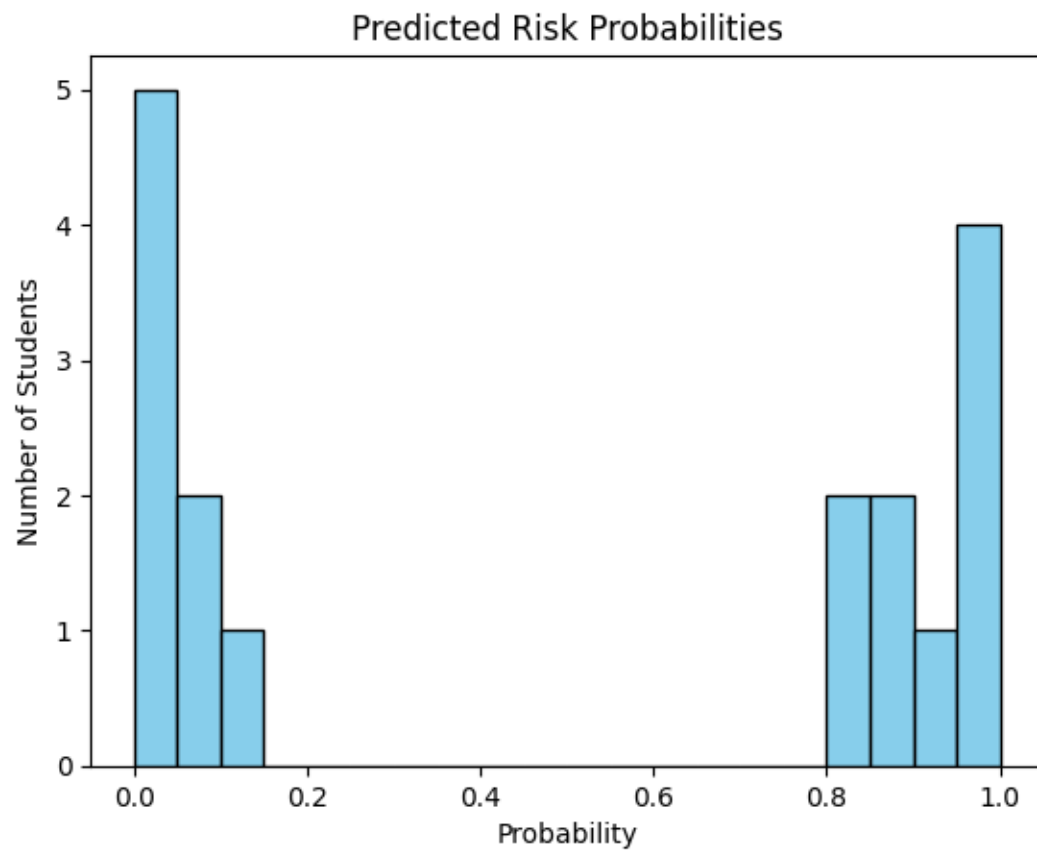
plt.ylabel('Risk Probability')

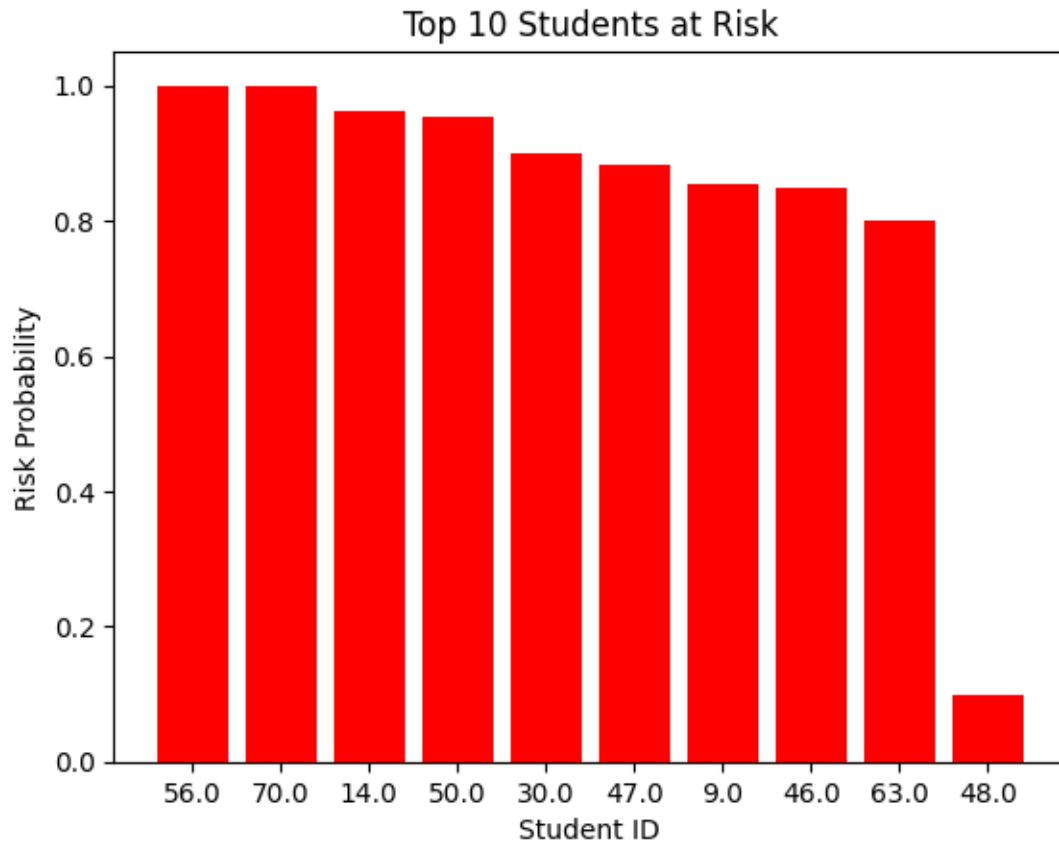
plt.title('Top 10 Students at Risk')

plt.show()
```

#### ----- STOP SPARK SESSION -----

```
spark.stop() print("☐ Spark session stopped. Project complete.")
```





## CHAPTER 10:

### Conclusion and Future Scope

The proposed system successfully predicts student risk levels using PySpark and Random Forest, demonstrating the potential of big data technologies in education. The use of distributed processing ensures scalability, while Random Forest ensures robust classification performance.

Future improvements could include integrating real-time data from learning management systems, applying advanced models like Gradient Boosted Trees or Neural Networks, and deploying the system as an interactive web application for educators.